



COMP 6721 Fall 2023 Project Part 2

Team Name: NS_13

Team Members: Wenhao Gu, Md Sayeed Abid, Archilkumar Dineshbhai Katrodiya

Student ID Numbers: 40203004, 40260779, 40270119

Team Member Specializations:

Data Specialist: Archilkumar Dineshbhai Katrodiya

Training Specialist: Wenhao Gu

Evaluation Specialist: Md Sayeed Abid

Project Repository Link: [[GitHub](#)]

Dataset Choices:

classes images	Image Source	Licensing Type	Relevant	Information
Neutral:	3522	FER-2013 Dataset	Publicly available	The dataset can be accessed at kaggle
Engaged/Focused:	3654	FER-2013 Dataset (choose by hand)	Publicly available	The dataset can be accessed at kaggle
Bored/Tired:	3714	FER-2013 Dataset (choose by hand)	Publicly available	The dataset can be accessed at kaggle
Angry/Irritated:	3744	FER-2013 Dataset	Publicly available	The dataset can be accessed at kaggle


We changed the dataset from part one because of an unbalanced reason.

In our part2, we focused on identifying 'engaged' and 'bored' emotions within the FER-2013 Dataset, which also includes Sad, Disgust, Fear, Surprise, and Happy categories. We handpicked approximately 700 images for each emotion (engaged and bored) from these categories. Our initial step was data cleaning, which involved carefully reviewing each image to ensure accurate categorization into 'engaged' or 'bored', as some images had ambiguous emotional expressions.

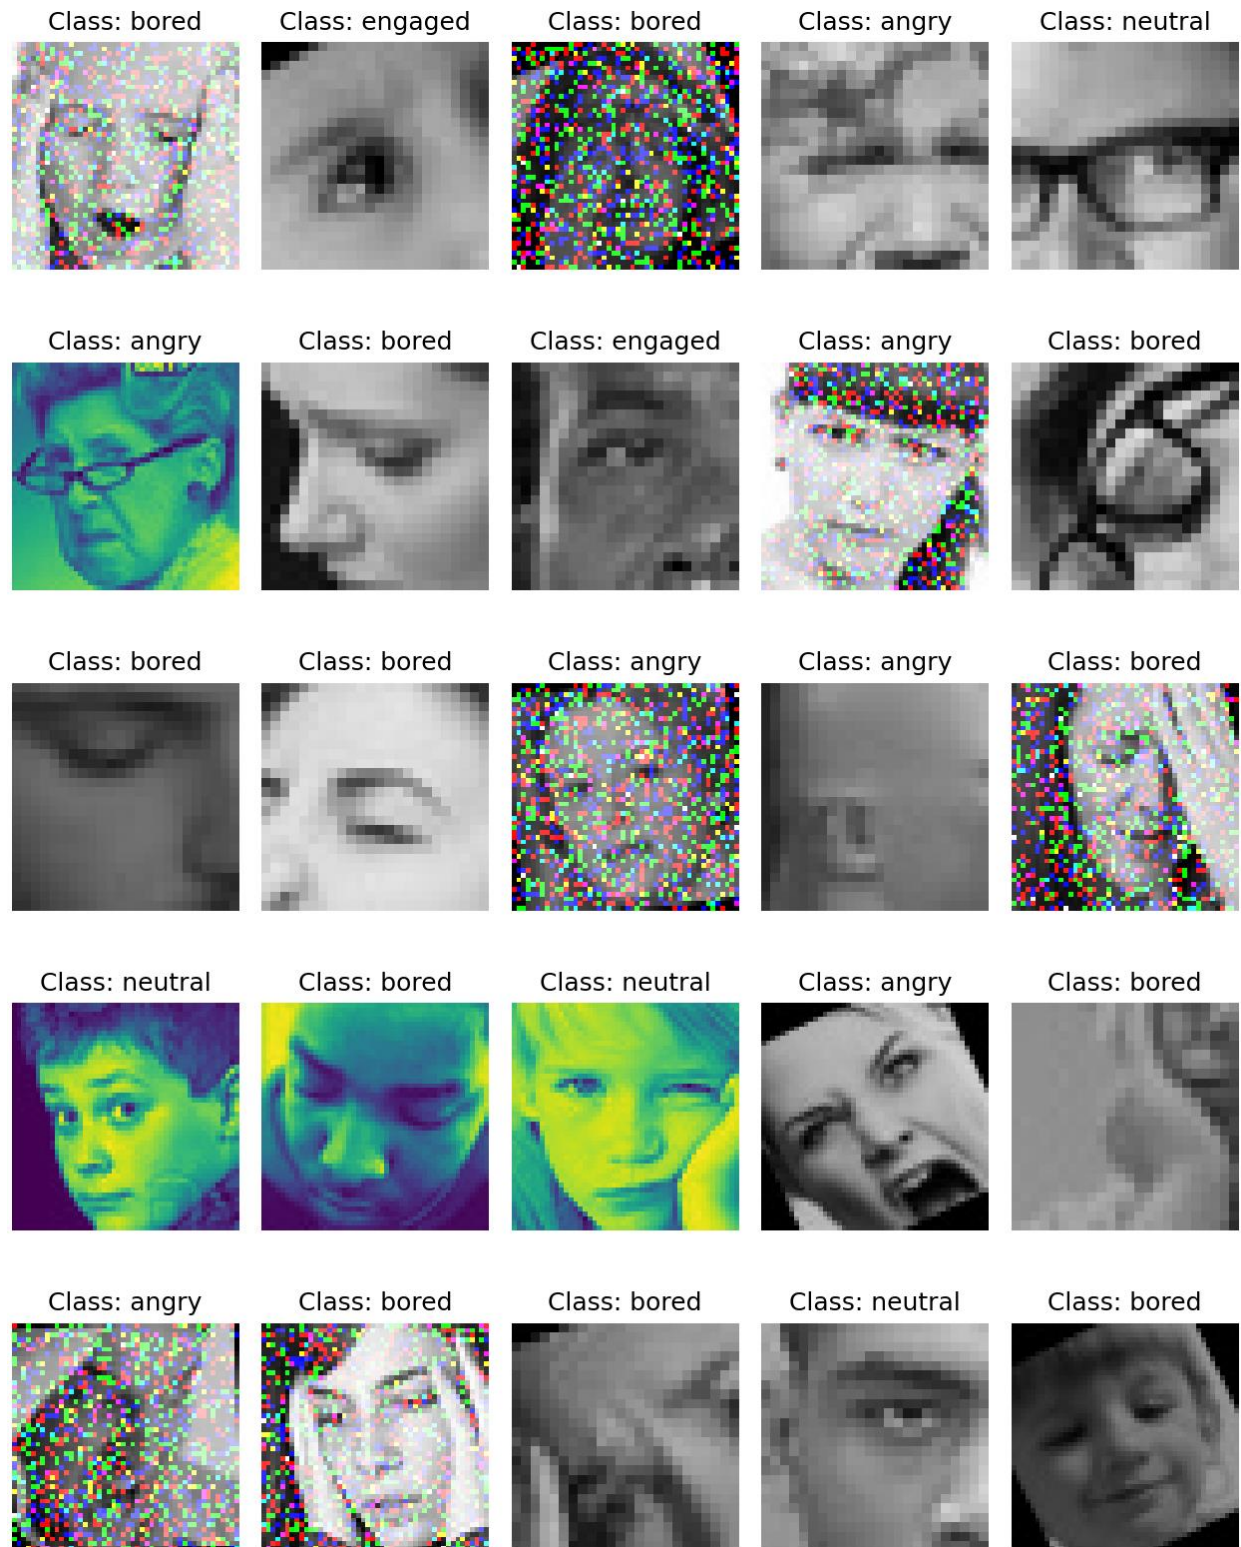
Subsequently, we enhanced the dataset through several image augmentation techniques. Each of the selected images underwent brightness correction and slight rotation. We also cropped the images for better focus on facial expressions and added noise to increase the robustness of our dataset. These augmentation techniques expanded our initial set of approximately 700 images per class to around 3600 images per class, enriching the dataset for more effective emotion recognition research.


Dataset Visualization:

Class Distribution:



Pixel Intensity Distribution:





CNN Architecture:

Main model:

The model has three convolutional layers for feature extraction and two fully connected layers for classification, 3*3 kernel size. We also added max-pooling and Dropout layers that enhance the model's generalization ability. We have used LeakyReLU as an activation function.

Convolution Layers:

Layer 1:

The first layer is the Convolution layer with kernel size 3 and padding 1. It produces output with 16 channels.

Layer 2:

Convolution layer with 16 input channels, 32 output channels and kernel size 2 with padding 1.

Layer 3:

Convolution layer with 32 input channels, 64 output channels and kernel size 3 with padding 1.

Pooling Layer:

We have used MaxPooling with Kernel size 2 and stride 2.

Fully Connected Layers:

The Hidden Layer contains 512 neurons, and the output layer contains 4.

Activation and Regularization:

- ReLU: Applied after each convolutional and fully connected layer except for the output layer.
- Dropout: Dropout with a dropout rate of 0.5 is imposed after the flattening and the first fully connected layer.

Model Variation 1:

Instead of three, this model has four convolution layers. In the first three convolution layers, input and output channels remain identical; however, kernel size is 3 for all the layers. Activation, pooling, fully connected, and dropout layers are the same. The invocation of activation and Dropout are also the same.

Convolution Layer 4:

Layer with 64 input channels, 128 output channels and kernel size 3 with padding 1.

Model variation 2:

The model has two convolution layers and only one fully connected layer. It uses LeakyReLU activation, dropout and max pooling for generalization.

Convolution Layers:

Layer 1:

The first layer is the Convolution layer with kernel size 3 with padding 1. It produces output with 16 channels.

Layer 2:

Convolution layer with 16 input channels, 32 output channels and kernel size 3 with padding 1.

Pooling Layer:

We have used MaxPooling with Kernel size 2 and stride 2.

Fully Connected Layers:

The output layer contains 4 neurons that predict the class.

Activation and Regularization:

- ReLU: Applied after each convolutional.
- Dropout: Dropout with a dropout rate of 0.5 is imposed after the flattening.

Training Process:

The models underwent training for 50 epochs using a learning rate of 0.001. We have chosen Cross-entropy loss as the loss function, while Adam is an optimizer.

We experimented with model training by varying the size of a kernel. We have used three different kernel sizes 2, 3 and 5. The model performed the best with a 3x3 kernel with 3 3-layer convolution net.

Results:

2x2 Kernel:

With the above kernel size, after 50 epochs, the model achieved a training accuracy of around 73.7% and a testing accuracy of 61.2%, 3 layers convolution net. It shows that the model experienced underfitting and did not perform as expected on the training or testing set. CNN is not able to extract enough information from the images. To overcome the problem, we increased the kernel size to 5.

Epoch: 48	Training Loss: 0.658150 0.914086	Training Accuracy: 0.731348 Validation Accuracy: 0.620404	Validation Loss:
	Validation loss decreased (0.921847 --> 0.914086). Saving model ...		
Epoch: 49	Training Loss: 0.659886 0.930901	Training Accuracy: 0.736035 Validation Accuracy: 0.612592	Validation Loss:
Epoch: 50	Training Loss: 0.658409 0.923049	Training Accuracy: 0.737402 Validation Accuracy: 0.620404	Validation Loss:
	Test Loss: 0.945618, Test Accuracy: 0.612592		

5x5 Kernel:

With the above kernel size, after 50 epochs, the model achieved a training accuracy of around 83.7% and a testing accuracy of 62.5%, 3 layers convolution net. The model was overfitting on the training set. To overcome the problem, we changed the kernel size to 3.

Epoch: 48	Training Loss: 0.441086 1.043358	Training Accuracy: 0.825000 Validation Accuracy: 0.638327	Validation Loss:
Epoch: 49	Training Loss: 0.425028 1.039613	Training Accuracy: 0.833984 Validation Accuracy: 0.630974	Validation Loss:
Epoch: 50	Training Loss: 0.419774 1.090913	Training Accuracy: 0.837402 Validation Accuracy: 0.625000	Validation Loss:
	Test Loss: 1.001778, Test Accuracy: 0.595588		

Main model 3x3 Kernel:

With the above kernel size, after 50 epochs, the model achieved a training accuracy of around 81% and a testing accuracy of 65%. It performed better compared to the other two kernels.

Epoch: 48	Training Loss: 0.491385 0.857007	Training Accuracy: 0.808008 Validation Accuracy: 0.665441	Validation Loss:
Epoch: 49	Training Loss: 0.503610 0.863198	Training Accuracy: 0.800293 Validation Accuracy: 0.665441	Validation Loss:
Epoch: 50	Training Loss: 0.486833 0.865220 Test Loss: 0.883289, Test Accuracy: 0.650276	Training Accuracy: 0.810059 Validation Accuracy: 0.657629	Validation Loss:

Variation 1:

The models change 3 layers convolution to 2 layers, under training for 50 epochs using a learning rate of 0.001 and a kernel size of 3. We have chosen Cross-entropy loss as the loss function, while Adam is an optimizer. The model achieved a training accuracy of around 57.7% and a testing accuracy of 50%. It shows that the model experienced underfitting.

Epoch: 48	Training Loss: 0.997496 1.125741	Training Accuracy: 0.570898 Validation Accuracy: 0.483915	Validation Loss:
Epoch: 49	Training Loss: 0.986199 1.104674	Training Accuracy: 0.574414 Validation Accuracy: 0.503676	Validation Loss:
Epoch: 50	Training Loss: 0.984567 1.108724 Test Loss: 1.104898, Test Accuracy: 0.509191	Training Accuracy: 0.577051 Validation Accuracy: 0.496783	Validation Loss:

Variation 2:

The models change 3 layers convolution to 4 layers, underwent training for 50 epochs using a learning rate of 0.001 and a kernel size of 3. We have chosen Cross-entropy loss as the loss function, while Adam is an optimizer. The model achieved a training accuracy of around 87% and a testing accuracy of 52.8%. It shows that the model experienced overfitting.

Epoch: 48	Training Loss: 0.326708 1.340406	Training Accuracy: 0.872559 Validation Accuracy: 0.581342	Validation Loss:
Epoch: 49	Training Loss: 0.332238 1.430754	Training Accuracy: 0.871680 Validation Accuracy: 0.592831	Validation Loss:
Epoch: 50	Training Loss: 0.338393 1.347836 Test Loss: 1.089433, Test Accuracy: 0.528493	Training Accuracy: 0.870898 Validation Accuracy: 0.585018	Validation Loss:

Evaluation:

Performance Metrics:


Model	Macro			Micro			Accuracy
	P	R	F1	P	R	F1	
Main Model	65.55	64.98	64.98	64.98	64.98	65.51	64.98
Variant1	50.81	50.72	50.87	50.87	50.87	50.81	50.87
Variant2	55.02	54.41	54.44	54.45	54.57	54.45	54.54

The main model has a training accuracy of 81%, and all the parameters are very close to it, which depicts that the model is generalized. Whereas, in variation 2, all the scores are low compared to its training accuracy of 90%, which shows that the model is overfitting on training data.


Confusion Matrix Analysis:

Based on the below confusion matrix, the most frequently confused class is natural. The main reason behind this misclassification is a tiny visual difference that increases false negatives. The model performs surprisingly well in recognizing angry and bored faces. Those classes have high visual differences compared to neutral and engaged, which create more discrete feature sets for the model to learn from and eventually reduce the model bias towards those classes.


Main model:



Variation 1:



Variation 2:



Impact of Architectural Variation:

We tried three different models. Variation 1 contains only two convolution layers. The model was unable to capture enough detail for classification. It underwent the problem of underfitting. We decided to increase the number of convolution layers to make the model extract detailed features. We increased the depth of the network by adding two more convolution layers. With 4 convolution layers, the model learned detailed features but lost generalization. We tried the third and our main model, in which we used three convolution layers. This model is capable of learning detailed features and is generalized.

We experimented with model training by varying the size of a kernel. We have used three different kernel sizes: 2, 3, and 5. The model performed the best with a 3x3 kernel. Detailed discussion about these experiments is described above in the **Training Process** part.

Conclusions and Forward Look:

The 3x3 Kernel with 3 Convolution Layers emerged as the best-performing model. It achieved a balance between training and testing accuracy (81% training, 65% testing), suggesting effective feature extraction without significant overfitting or underfitting. Smaller kernels (2x2) were insufficient for feature extraction and lacked the model complexity needed to capture essential features, while larger ones (5x5) led to overfitting. More layers (4 layers) led to overfitting, whereas fewer layers (2 layers) resulted in underfitting. A 3-layer architecture appeared optimal for your dataset. Each change in the model configuration had a significant impact on performance.

Consider employing data augmentation approaches to further increase generalization. This can include rotating, resizing, or flipping photos to provide a more diversified training set, which will help the model generalize better. Consider pre-trained models and fine-tuning them on your specific data for datasets similar to typical image datasets. This technique can make use of learned features from big, heterogeneous datasets. Dropout or L2 regularization should be used to prevent overfitting, especially in models with increased complexity (such as the 4-layer form).